


Qwen2.5-7B-Instruct (4-bit): accuracy vs budget



Qwen2.5-7B-Instruct (4-bit): compute vs budget

